

FACULDADE DE ENGENHARIA DA UNIVERSIDADE DO PORTO

Android-based implementation of Eulerian Video Magnification for vital signs monitoring

Pedro Boloto Chambino

WORKING VERSION



Mestrado Integrado em Engenharia Informática e Computação

Supervisor: Luís Teixeira

Supervisor at Fraunhofer Portugal: Luís Rosado

May 30, 2013

Android-based implementation of Eulerian Video Magnification for vital signs monitoring

Pedro Boloto Chambino

Mestrado Integrado em Engenharia Informática e Computação

May 30, 2013

Abstract

Eulerian Video Magnification is a recently presented method capable of revealing temporal variations in videos that are impossible to see with the naked eye. Using this method, it is possible to visualize the flow of blood as it fills the face. From its result, a person's heart rate is possible to be extracted.

This research work is an internal project of *Fraunhofer Portugal* and its goal is to test the feasibility of the implementation of the Eulerian Video Magnification method on smartphones by developing an *Android* application for monitoring vital signs based on the Eulerian Video Magnification method.

There has been some successful effort on the assessment of vital signs, such as, heart rate, and breathing rate, in a contact-free way using a webcam and even a smartphone. However, since the Eulerian Video Magnification method was recently proposed, its implementation has not been tested in smartphones yet.

The application will include features, such as, detection of a person's cardiac pulse, dealing with artifacts' motion, and real-time display of the magnified blood flow. Then, the application performance will be evaluated through tests with several individuals and the assessed heart rate compared to the one detected by the *Philips* application, and to the measurement of an heart rate monitor or a pulse oximeter.

Contents

1	Introduction	1
1.1	Context	1
1.2	Motivation	2
1.3	Objectives	2
1.4	Outline	3
2	State of the art	5
2.1	Photo-plethysmography	5
2.2	Signal post-processing	6
2.2.1	Independent Component Analysis	6
2.2.2	Eulerian Video Magnification	7
2.2.3	Detrending	10
2.3	Heart rate estimation	10
2.3.1	Power spectrum	10
2.3.2	Pulse wave detection	11
2.4	Technologies	12
2.4.1	Android SDK	12
2.4.2	OpenCV – Computer Vision Library	13
2.5	Chapter summary	13
3	Problem description	15
3.1	Android-based implementation of Eulerian Video Magnification	15
3.2	Vital signs monitoring	16
3.3	Chapter summary	16
4	Implementation details	17
4.1	Overview	17
4.2	Eulerian Video Magnification implementations	18
4.2.1	EvmGdownIdeal	19
4.2.2	EvmGdownIIR	20
4.2.3	EvmLpyrIIR	20
4.2.4	Performance optimized EvmGdownIIR	21
4.3	Face detection	22
4.4	Signal validations	23
4.5	Heart rate estimation	23
4.6	Performance optimizations	23
4.7	Android integration	23
4.8	Chapter summary	23

CONTENTS

5	Results	25
5.1	Heart rate comparison	25
5.2	Algorithm performance	25
5.3	Chapter summary	25
6	Conclusions	27
6.1	Objective satisfaction	27
6.2	Future work	27
	References	29

List of Figures

2.1	Overview of the Eulerian Video Magnification method.	7
2.2	Examples of temporal filters.	8
2.3	Emphasis of face color changes using the Eulerian Video Magnification method. .	9
2.4	Original and detrended RR series.	10
4.1	Implemented library modules overview.	18
4.2	Eulerian Video Magnification method steps.	19
4.3	Overview of image deconstruction and reconstruction for building a Laplacian pyramid.	20
4.4	Integration workflow between Android native and Java parts.	23

LIST OF FIGURES

List of Tables

LIST OF TABLES

Abbreviations

EVM	Eulerian Video Magnification
ICA	Independent Component Analysis
PPG	Photo-plethysmography
FFT	Fast Fourier transform
FPS	Frames per second
JNI	Java Native Interface
JVM	Java Virtual Machine
OpenCV	Open Source Computer Vision Library
IIR	Infinite impulse response

Chapter 1

Introduction

This chapter introduces this work, by first presenting its context, motivation, and project’s objectives, on sections 1.1, 1.2, and 1.3, respectively.

Finally, section 1.4 describes the document outline.

1.1 Context

Eulerian Video Magnification is a method, recently presented at *SIGGRAPH*¹ 2012, capable of revealing temporal variations in videos that are impossible to see with the naked eye. Using this method, it is possible to visualize the flow of blood as it fills the face [WRS⁺12]. Which provides enough information to assess the heart rate in a contact-free way using a camera [WRS⁺12, PMP10, PMP11].

The main field of this research work is *image processing and computer vision*, whose main purpose is to translate dimensional data from the real world in the form of images into numerical or symbolical information.

Other fields include *medical applications, software development for mobile devices, digital signal processing*.

This research work is an internal project of *Fraunhofer Portugal*² supervised by Luís Rosado. Fraunhofer Portugal a is non-profit private association founded by Fraunhofer-Gesellschaft³ [Por13] and

“aims on the creation of scientific knowledge capable of generating added value to its clients and partners, exploring technology innovations oriented towards economic growth, the social well-being and the improvement of the quality of life of its end-users.” [Por13]

¹<http://www.siggraph.org/>

²<http://www.fraunhofer.pt/>

³<http://www.fraunhofer.de/en/about-fraunhofer/>

1.2 Motivation

Due to being recently proposed, the Eulerian Video Magnification method implementation has not been tested in smartphones yet.

There has been some successful effort on the assessment of vital signs, such as, heart rate, and breathing rate, in a contact-free way using a webcam [WRS⁺12, PMP10, PMP11], and even a smartphone [Tec13, Phi13].

Other similar products, which require specialist hardware and are thus expensive, include *laser Doppler* [UT93], *microwave Doppler radar* [Gre97], and *thermal imaging* [GSMP07].

Since it is a cheaper method of assessing vital signs in a contact-free way than the above products, this research work has potential for advancing fields, such as, *telemedicine*, *personal health-care*, and *ambient assisting living*.

Despite the existence of very similar products by *Philips* [Phi13] and *ViTrox Technologies* [Tec13] to the one proposed on this research work, none of these implements the Eulerian Video Magnification method. Moreover, the application to be developed during this research work will have additional features described on the next section 1.3.

1.3 Objectives

This research work goal is to test the feasibility of the implementation of the Eulerian Video Magnification method on smartphones by developing an *Android* application for monitoring vital signs based on the Eulerian Video Magnification method.

This application should include the following features:

- heart rate detection and assessment based on the Eulerian Video Magnification method;
- display real-time changes, such as, the magnified blood flow, obtained from the Eulerian Video Magnification method;
- deal with artifacts' motion, due to, person and/or smartphone movement.

It should be noted that a straightforward implementation of the Eulerian Video Magnification method is not possible, due to various reasons. First, the Eulerian Video Magnification method provides motion magnification along with color magnification which will introduce several problems with artifacts' motion. Second, the requirement of implementing a real-time smartphone application will create performance issues which will have to be addressed and trade-offs will have to be considered.

The application performance should then be evaluated through tests with several individuals and the assessed heart rate compared to the ones detected by another application [Tec13, Phi13], and to the measurement of an electronic sphygmomanometer.

1.4 Outline

The rest of the document is structured as follows:

Chapter 2 introduces the concepts necessary to understand the presented problem. In addition, it presents the existing related work, and a description of the technologies to be used.

Chapter 3 ...

Chapter ?? presents the approach taken to solve the problem. Moreover, it introduces the testing and evaluation methodologies.

Chapter 4 ...

Chapter 5 ...

Chapter 6 ...

Introduction

Chapter 2

State of the art

This chapter presents several studies regarding the heart rate estimation from a person's face captured through a simple webcam.

Section 2.1 describe the concept that explains how the cardiac pulse is detected from a person's face in a remote, contact-free way.

Post-processing methods, which may be applied to the retrieved signal, are detailed on section 2.2.

In order to estimate the heart rate, a couple of techniques are also detailed on section 2.3.

Finally, section 2.4 reviews the main technologies and tools used throughout this work.

2.1 Photo-plethysmography

Photo-plethysmography (PPG) is the concept of measuring volumetric changes of an organ optically. Its most established use is in pulse oximeters.

PPG is based on the principle that blood absorbs more light than surrounding tissue thus variations on blood volume affect light reflectance [VSN08].

The use of dedicated light sources and infra-red wavelengths, and contact probes has been the norm [UT93, Gre97, GSMP07]. However, recently, remote, non-contact PPG imaging has been explored.

The method used on the article [VSN08] captures the pixel values (red, green, and blue channels) of the facial area of a previously recorded video where volunteers were asked to minimize movements. The pixel values within a region of interest (ROI) was then averaged for each frame. This spatial averaging was found to significantly increase signal-to-noise ratio. The heart rate estimation was then calculated by applying Fast Fourier transforms and the power spectrum as explained on section 2.3.1.

The authors of [VSN08] demonstrate that the fact that the green channel features a stronger heart rate signal as compared to the red and blue channels, is a strong evidence that the signal is due to variations in blood volume because (oxy-) hemoglobin absorbs green light.

2.2 Signal post-processing

After obtaining the raw pixel values (red, green, and blue channels), a conjunction of the following methods may be used to extract and improve the reflected plethysmography signal. However, each method introduces complexity and expensive computation.

2.2.1 Independent Component Analysis

Independent Component Analysis is a special case of *blind source separation* and is a relatively new technique for uncovering independent signals from a set of observations that are composed of linear mixtures of the underlying sources [Com94].

In this case, the underlying source signal of interest is the cardiac pulse that propagates throughout the body, which modify the path length of the incident ambient light due to volumetric changes in the facial blood vessels during the cardiac cycle, such that subsequent changes in amount of reflected light indicate the timing of cardiovascular events.

By recording a video of the facial region, the red, green, and blue (RGB) color sensors pick up a mixture of the reflected plethysmographic signal along with other sources of fluctuations in light due to artifacts. Each color sensor records a mixture of the original source signals with slightly different weights. These observed signals from the red, green and blue color sensors are denoted by $x_1(t)$, $x_2(t)$ and $x_3(t)$ respectively, which are amplitudes of the recorded signals at time point t . In conventional Independent Component Analysis model the number of recoverable sources cannot exceed the number of observations, thus three underlying source signals were assumed, represented by $s_1(t)$, $s_2(t)$ and $s_3(t)$. The Independent Component Analysis model assumes that the observed signals are linear mixtures of the sources, i.e. $x_i(t) = \sum_{j=1}^3 a_{ij}s_j(t)$ for each $i = 1, 2, 3$. This can be represented compactly by the mixing equation

$$x(t) = As(t) \quad (2.1)$$

where the column vectors $x(t) = [x_1(t), x_2(t), x_3(t)]^T$, $s(t) = [s_1(t), s_2(t), s_3(t)]^T$ and the square 3×3 matrix A contains the mixture coefficients a_{ij} . The aim of Independent Component Analysis model is to find a separating or demixing matrix W that is an approximation of the inverse of the original mixing matrix A whose output

$$\hat{s}(t) = Wx(t) \quad (2.2)$$



Figure 2.1: Overview of the Eulerian Video Magnification method.

is an estimate of the vector $s(t)$ containing the underlying source signals. To uncover the independent sources, W must maximize the non-Gaussianity of each source. In practice, iterative methods are used to maximize or minimize a given cost function that measures non-Gaussianity [PMP10, PMP11].

2.2.2 Eulerian Video Magnification

In contrast to the Independent Component Analysis model that focus on extracting a single number, the Eulerian Video Magnification uses localized spatial pooling and temporal filtering to extract and reveal visually the signal corresponding to the cardiac pulse. This allows for amplification and visualization of the heart rate signal at each location on the face. This creates potential for monitoring and diagnostic applications to medicine, i.e. the asymmetry in facial blood flow can be a symptom of arterial problems.

Besides color amplification, the Eulerian Video Magnification method is also able to reveal low-amplitude motion which may be hard or impossible for humans to see. Previous attempts to unveil imperceptible motions in videos have been made, such as, [LTF⁺05] which follows a *Lagrangian* perspective, as in fluid dynamics where the trajectory of particles is tracked over time. By relying on accurate motion estimation and additional techniques to produce good quality synthesis, such as, motion segmentation and image in-painting, the algorithm complexity and computation is expensive and difficult.

On the contrary, the Eulerian Video Magnification method is inspired by the *Eulerian* perspective, where properties of a voxel of fluid, such as pressure and velocity, evolve over time. The approach of this method to motion magnification is the exaggeration of motion by amplifying temporal color changes at fixed positions, instead of, explicitly estimation of motion.

This method approach, illustrated in figure 2.1, combines spatial and temporal processing to emphasize subtle temporal changes in a video. First, the video sequence is decomposed into different spatial frequency bands. Because they may exhibit different signal-to-noise ratios, they may be magnified differently. In the general case, the full Laplacian pyramid [BA83] may be



Figure 2.2: Examples of temporal filters.

computed. Then, temporal processing is performed on each spatial band. The temporal processing is uniform for all spatial bands, and for all pixels within each band. After that, the extracted bandpass signal is magnified by a factor of α , which can be specified by the user, and may be attenuated automatically. Finally, the magnified signal is added to the original and the spatial pyramid collapsed to obtain the final output.

2.2.2.1 Spatial filtering

As mention before, the work of [WRS⁺12] computes the full Laplacian pyramid [BA83] as a general case for spatial filtering. Each layer of the pyramid may be magnified differently because it may exhibit different signal-to-noise ratios, or contain spatial frequencies for which the linear approximation used in motion magnification does not hold [WRS⁺12, Section 3].

Spatial filtering may also be used to significantly increases signal-to-noise ratio, as previously mention on section 2.1 and demonstrated on the work of [VSN08] and [WRS⁺12]. Subtle signals, such as, a person’s heart rate from a video of its face, may be enhanced this way. For this purpose the work of [WRS⁺12] computes a layer of the Gaussian pyramid which may be obtained by successively scaling down the image by calculating the Gaussian average for each pixel.

However, for the signal of interest to be revealed, the spatial filter applied must be large enough. Section 5 of [WRS⁺12] provides an equation to estimate the size for a spatial filter needed to reveal a signal at a certain noise power level:

$$S(\lambda) = S(r) = \sigma'^2 = k \frac{\sigma^2}{r^2} \quad (2.3)$$

where $S(\lambda)$ represents the signal over spatial frequencies, and since the wavelength, λ , cutoff of a spatial filter is proportional to its radius, r , the signal may be represented as $S(r)$. The noise power, σ^2 , can be estimated using to the technique of [LFSK06]. Finally, because the filtered noise power level, σ'^2 , is inversely proportional to r^2 , it is possible to solve the equation for r , where k is a constant that depends on the shape of the low pass filter.

2.2.2.2 Temporal filtering

Temporal filtering is used to extract the motions or signals to be amplified. Thus, the filter choice is application dependent. For motion magnification, a broad bandpass filter, such as, the butterworth filter, is preferred. A narrow bandpass filter produces a more noise-free result for color



Figure 2.3: Emphasis of face color changes using the Eulerian Video Magnification method.

amplification of blood flow. An ideal bandpass filter is used on [WRS⁺12] due to its sharp cutoff frequencies. Alternatively, for a real-time implementation low-order IIR filters can be useful for both: color amplification and motion magnification. These filters are illustrated on 2.2.

2.2.2.3 Emphasize color variations for human pulse

The extraction of a person's cardiac pulse using the Eulerian Video Magnification method was demonstrated in [WRS⁺12]. It was also presented that using the right configuration can help extract the desired signal. There are four steps to take when processing a video using the Eulerian Video Magnification method:

1. select a temporal bandpass filter;
2. select an amplification factor, α ;
3. select a spatial frequency cutoff (specified by spatial wavelength, λ_c) beyond which an attenuated version of α is used;
4. select the form of the attenuation for α — either force α to zero for all $\lambda < \lambda_c$, or linearly scale α down to zero.

For human pulse color variation, two temporal filters may be used, first selecting frequencies within 0.4-4Hz, corresponding to 24-240 beats per minute (bpm), then a narrow band of 0.83-1Hz (50-60 bpm) may be used, if the extraction of the pulse rate was successful.

To emphasize the color change as much as possible, a large amplification factor, $\alpha \approx 100$, and spatial frequency cutoff, $\lambda_c \approx 1000$, is applied. With an attenuation of α to zero for spatial wavelengths below λ_c .

The resulting output can be seen in figure 2.3.



Figure 2.4: Original and detrended RR series.

2.2.3 Detrending

Detrending is a method of removing very large ultralow-frequency trends an input signal without any magnitude distortion, acting as an high-pass filter.

The main advantage of the method presented on the work of [TRaK02], compared to methods presented in [LOCS95] and [PB90], is its simplicity.

The method consists of separating the input signal, z , into two components, as $z = z_{stat} + z_{trend}$, where z_{stat} is the nearly stationary component, and z_{trend} is the low frequency aperiodic trend component.

An estimation of the nearly stationary component, \hat{z}_{stat} , can be obtained using the equation below. The detailed derivation of the equation can be found in [TRaK02].

$$\hat{z}_{stat} = (I - (I + \lambda^2 D_2^T D_2)^{-1})z \quad (2.4)$$

where I is the identity matrix, D_2 is the discrete approximation of the second order, and λ is the regularization parameter.

Figure 2.4 presents an example of what this method is able to achieve. The example, taken from the work of [TRaK02], uses real RR series and the effect of the method on time and frequency domain analysis of heart rate variability is demonstrated not to lose any useful information.

2.3 Heart rate estimation

In order to convert the extracted plethysmographic signal into the number of beats per minute (bpm), further processing must be done. Below are highlighted two methods capable of achieving this goal.

2.3.1 Power spectrum

Fourier transform is a mathematical transform capable of converting a function of time, $f(t)$, into a new function representing the frequency domain of the original function.

To calculate the power spectrum, the resulting function from the *Fourier transform* is then multiplied by itself.

Since the values are captured from a video, sequence of frames, the function of time is actually discrete, with a frequency rate equal to the video frame rate, FPS .

The *index*, i , corresponding to the maximum of the power spectrum can then be converted into a frequency value, F , using the equation:

$$F = \frac{i * FPS}{2N} \quad (2.5)$$

where N is the size of the signal extracted. F can then be multiplied by 60 to convert it to beats per minute, and have an estimation of the heart rate from the extracted signal.

2.3.2 Pulse wave detection

In [NI10], it is presented an automated algorithm for fast pulse wave detection. The algorithm is capable of obtaining an estimative of the heart rate from PPG signal, as an alternative to the power spectrum described above. Moreover, it also introduces validation to the waveform detection by verifying its shape and timing. Below is presented a simplified description of the algorithm. A more detailed description can be found in [NI10].

1. Identification of possible peaks and foots of individual pulses

(a) Maximum (*MAX*)

The signal is divided into consecutive 200ms time intervals and for every segment the absolute maximum is determined. Some of these maximums are rejected: if they fall below a predetermined amplitude threshold; or if the distance between two maximums is less than or equal to 200ms, then the lower maximum is rejected.

(b) Minimum (*MIN*)

The absolute minimum is determined between every two adjacent maximums. A minimum is rejected, it is above a predetermined amplitude threshold. When a minimum is rejected, the lower-amplitude maximum of the two maximum adjacent to the rejected minimum is discarded too.

2. Examination and verification of the rising edges

(a) Validation of a single rising edge

If a rising edge is rejected, its maximum and minimum are rejected. A rising edge is rejected, if its amplitude ($AMPL = MAX - MIN$) is lower than amplitude threshold; or its duration is lower than a threshold that depends on the sampling rate; or its amplitude does not increase smoothly.

(b) Estimation of the similarity of a rising edge to preceding and following rising edges accepted as valid

Two rising edges are considered similar, if the amplitude of the lower-amplitude rising edge is greater than 50% of the amplitude of the higher-amplitude rising edge; and if the maximum of the lower-amplitude rising edge is between $\pm 60\%$ of the maximum of the higher-amplitude rising edge; and if the minimum of the lower-amplitude rising edge is between $\pm 60\%$ of the minimum of the higher-amplitude rising edge; and if the

duration of the shorter rising edge is greater than 33% of the duration of the longer rising edge. The valid rising edges are then categorized according to its characteristics for the following step. The categorization description is suppressed for brevity and can be found at [NI10].

(c) Verification of the current rising edge

The rising edges categorized on the previous step are considered valid edges of a pulse wave if they fulfill at least one of the decision rules presented on [NI10] and suppressed for brevity.

The validation process described here is important for discarding signals which are not representative of pulse waves. Providing a way of calculating the heart rate estimation only on valid pulse signals.

2.4 Technologies

Below are short descriptions of two of the main technologies that will be used during this research work.

2.4.1 Android SDK

Android SDK is the development kit for the *Android* platform. The *Android* platform is an open source, Linux-based operating system, primarily designed for touchscreen mobile devices, such as, smartphones.

Because of its open source code and permissive licensing, it allows the software to be freely modified and distributed. This have allowed *Android* to be the software of choice for technology companies who require a low-cost, customizable, and lightweight operating system for mobile devices and others.

Android has also become the world's most widely used smartphone platform with a worldwide smartphone market share of 75% during the third quarter of 2012 [IDC13].

Android consists of a kernel based on Linux kernel with middleware, libraries and APIs written in C. Applications, usually, run on an application framework which includes Java-compatible libraries based on *Apache Harmony*, an open source, free Java implementation. *Java bytecode* is then translated to run on the *Dalvik virtual machine*.

Porting existing Linux application or libraries to *Android* is difficult due to the lack of a native *X Window System* and lack of support for *GNU* libraries. Support for simple C and SDL application is possible, though, by the usage of *JNI*, a programming framework that allows Java code to call and be called by libraries written in C/C++.

2.4.2 OpenCV – Computer Vision Library

OpenCV is a library of programming functions mainly aimed at real-time image processing. To support these, it also includes a statistical machine learning library. Moreover, it is a cross-platform and open source library that is free to use and modify under the BSD license.

“OpenCV was built to provide a common infrastructure for computer vision applications and to accelerate the use of machine perception in the commercial products.” [Its13]

OpenCV is written in C/C++. There are binding for other languages, such as, Python, Java, and even Android. However, Java and Android implementation is recent and lacks features and stability.

2.5 Chapter summary

This chapter starts by describing the concept behind the extraction of cardiac pulse is possible from a person’s face captured through a simple video or webcam.

It then presents several possible post-processing methods for to improve the extraction of the actual pulse signal. These methods include:

- *Independent Component Analysis*, a method capable of uncovering independent signals from a set of observations that are composed of linear mixtures of the underlying sources;
- *Eulerian Video Magnification*, a method inspired by the *Eulerian* perspective that exaggerates color variations by analyzing how each pixel value changes over time;
- *Detrend*, a method which removes small trends from an input signal without distorting its amplitude.

Then algorithms for obtaining the actual beats per minutes of the heart rate from the signal are described:

- *Power spectrum*, a set of equations capable of finding the frequency of a signal using the *Fourier transform*;
- *Pulse wave detection*, an algorithm for detecting and validating rising edges from a pulse signal.

Finally, important technologies for the work are described and explored:

- *Android*, a Linux-based operating system, primarily designed for touchscreen mobile devices;
- *OpenCV*, a *Computer Vision* library of programming functions mainly aimed at real-time image processing.

State of the art

Chapter 3

Problem description

This chapter provides a detailed description of the problem addressed, defining its scope and dividing it in smaller problems.

Section 3.1 describes the main objective of the work which consists of an implementation a video magnification method based on the Eulerian perspective capable of running on a mobile device.

Then, section 3.2 provides a description of a simple application of the Eulerian Video Magnification method.

3.1 Android-based implementation of Eulerian Video Magnification

Fraunhofer Portugal is interested in testing the feasibility of implementing an Eulerian Video Magnification-based method on a mobile device with the Android platform.

As stated on the previous chapters, the Eulerian Video Magnification method is capable of magnifying small motion and amplifying color variation which may be invisible to the naked eye. Examples of the method application include: estimation of a person's heart rate from the variation of its face's color; respiratory rate from a person's chest movements; and even, detect asymmetry in facial blood flow, which may be a symptom of arterial problems.

The benefits of the Eulerian perspective is its low requirements for computational resources and algorithm complexity, in comparison to other attempts which rely on accurate motion estimation [LTF⁺05]. However, the existing limits of computational power on mobile devices may not allow the Eulerian Video Magnification method to execute in real-time.

The main project's goal is to develop a lightweight, real-time Eulerian Video Magnification-based method capable of executing on a mobile device. Which will require performance optimizations and trade-offs will have to taken into account.

3.2 Vital signs monitoring

As an objective to demonstrate that the Eulerian Video Magnification-based method developed is working as expected, the creation of an Android application which estimates a person's heart rate in real-time using the device's camera was pursued.

This goal requires comprehension of the photo-plethysmography concept, extraction of a frequency from a signal, and recognition / validation of a signal as a cardiac pulse.

The application will then need to be tested in order to verify its estimations. The test will be achieved by comparing results from a sphygmomanometer and other existing application [[Tec13](#), [Phi13](#)] which use different methods to estimate a person's heart rate.

3.3 Chapter summary

In this chapter, a more detailed description of the problem and its scope is presented.

It describes the main goal and motivation for developing a lightweight, real-time Eulerian Video Magnification-based method for the Android platform.

Moreover, the goal of creating an Android application for heart rate monitoring is explained. Which serves for testing the developed method and demonstrating what can be achieved with the implemented Eulerian Video Magnification method.

Chapter 4

Implementation details

This chapter ...

4.1 Overview

In order to create an Android application capable of estimating a person's heart rate, a desktop application was developed because of its faster implementation speed, and easier testing. The application was divided into several modules, illustrated in figure 4.1, which, later, were extracted into an independent library to be integrated into an Android application. The language used to implement the desktop application and library were C/C++. In addition, for the image processing operations, the computer vision library, OpenCV, was used.

A short description of the modules on figure 4.1 and application workflow is as follows:

1. *Original frame*, read frame from device's webcam;
2. *Face detector*, detect faces in current frame and match with previously detected faces in order to track multiple faces. Each face information is then fed into the following modules:
 - (a) *Eulerian Video Magnification*, magnify detected face's rectangle;
 - (b) *Signal noise*, verify if the signal extracted from the current face is too noisy. If so, that face's signal is reset and marked as not valid;
 - (c) *Detrending & Normalization*, if the current face's signal is not too noisy, then detrend and normalize the signal in order to facilitate further operations with the signal;
 - (d) *Validate signal*, the face's signal is then validated by verifying its shape and timing, in a similar but simpler manner as found in [NI10]. If the signal is given as invalid, it is kept as valid for a couple of time, because the validation algorithm may miss some peaks;

Implementation details

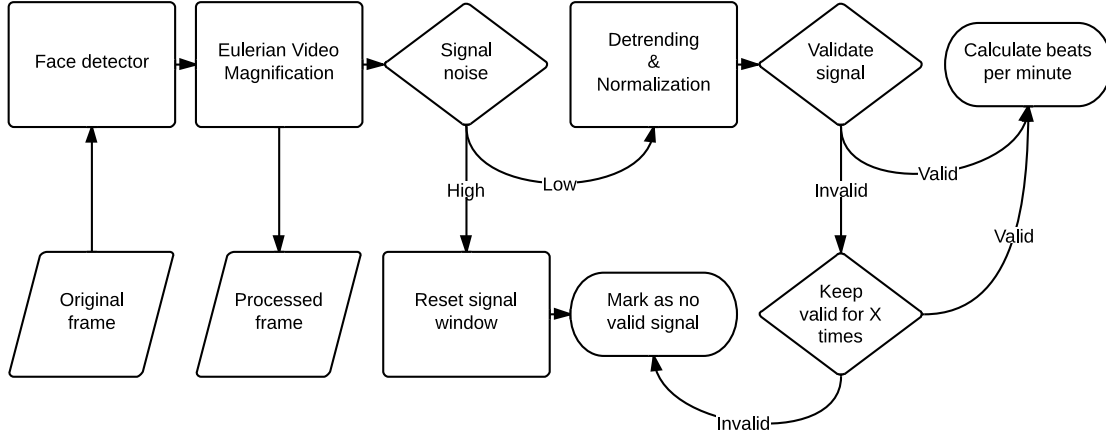


Figure 4.1: Implemented library modules overview.

(e) *Calculate beats per minute*, if the current face's signal is valid, it is then used to estimate the person's heart rate;

3. *Processed frame*, the resulting frame with each magnified face rectangle added back to the original frame.

Below, there are more detailed information and descriptions of the modules.

4.2 Eulerian Video Magnification implementations

This section presents the details of several different implementations of the Eulerian Video Magnification method.

The first implementations, described on sections 4.2.1, 4.2.2 and 4.2.3, were developed in Java to facilitate the integration into the Android application. However, the OpenCV Java binding was still in its early stages which end up creating difficulties for the development. Thus, the final implementation, on section 4.2.4, was implemented in C/C++, which also reduces the number of JNI calls from the Android JVM and increases the application performance.

The purpose of implementing multiple variants of the method was to study how the method worked and select which spatial and temporal filters would better fit the application goal: amplify color variation in real-time.

Figure 4.2 shows generic steps of the method which will be detailed on each of the following sections. The final step, *add to original frame*, however, remains the same in all implementations. Which is when the magnified values are added back to the original frame in order to obtain the processed frame.




Figure 4.2: Eulerian Video Magnification method steps.

4.2.1 EvmGdownIdeal

This was the first implementation, thus, its goal was to understand how the method worked, and match the implementation provided, in MATLAB, by [WRS⁺12]. However, this implementation would have real-time support by using a sliding window of 30 frames.

Resize down

This step applies a spatial filter by calculating a level of the Gaussian pyramid. This is achieved by looping to the desired level where the input to the next loop is the result from the previous loop, starting with the original frame. A Gaussian pyramid level is calculated by, first, convolving the input frame with the kernel:

$$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix} \quad (4.1)$$

and then, downsampling the frame by rejecting even rows and columns.

Temporal filter

It was used an ideal bandpass filter to remove any amplification of undesired frequency from the color variation of each pixel. To construct this ideal filter, the Fourier transform was calculated for each pixel over the sliding window of 30 frames. Then, frequencies below 45 and above 240 were set to zero, and the frame was rebuilt using the inverse Fourier transform.

Amplification

In this step, the result of the temporal filter is multiplied by an α value. Which results in the magnification of the color variation selected by the temporal filter.

Implementation details




Figure 4.3: Overview of image deconstruction and reconstruction for building a Laplacian pyramid.

Resize up

This step performs the inverse operation of the *resize down* step, where it upsamples the frame by inserting even rows and columns with zeros, and then, convolves the input frame with the same kernel multiplied by 4. However, when the original frame is not multiple of two, an additional resize operation as to be done in order for the upsampled frame to match the original frame's size.

4.2.2 EvmGdownIIR

This implementation is very similar to the one above, but uses a different temporal filter which does not require a sliding window of frames to support real-time results. The filter used was an IIR bandpass filter, which was constructed from the subtraction of two first-order lowpass IIR filters. Each lowpass filter is computed as follows:

$$L_n = L_{n-1} * (1 - \omega) + \omega * M \quad (4.2)$$

where M is the current frame, L is the lowpass filter accumulator for each frame, and ω is the cutoff frequency percentage.

The IIR temporal bandpass filter demonstrated similar results to the ideal temporal filter used on the first implementation, without the need for persisting a sliding window of frames, which simplifies solution and reduces the computational power required by the device.

4.2.3 EvmLpyrIIR

Using the same IIR temporal filter as above, this implementation uses a different spatial filter, which, instead of, computing a level of the Gaussian pyramid, it constructs the full Laplacian pyramid and then applies the temporal filter to each of its bands and each band is amplified differently.

Resize down

Figure 4.3 shows the steps to decompose and reconstruct an image for the purpose of building a Laplacian pyramid. The *original image* must be decomposed into two images, *blurred* and *fine*, by applying any type of spatial lowpass filter and scaling the image down or up by 2. In this case, a Gaussian filter was applied as described on steps *Resize down* and *Resize up* of the first implementation. Further levels of the pyramid can be computed by decomposing the *blurred image* in the same manner.

Temporal filter

The temporal filter used is the IIR bandpass filter, as described above for the previous implementation, only this time it is applied to each level of the pyramid.

Amplification

The amplification method in this implementation is more complex than the one previously used. It is based on the implementation provided by [WRS⁺12]. It uses a different α value for each band of spatial frequencies, which corresponds to the Laplacian pyramid levels. The magnification value, α , follows the equation:

$$(1 + \alpha)\delta(t) < \frac{\lambda}{8} \quad (4.3)$$

where $\delta(t)$ represents the displacement function and λ the spatial wavelength. Further details about this equation may be found on [WRS⁺12, Section 3.2].

Resize up

This step reconstructs the *original image* by iteratively reconstructing each *blurred image* until the now processed *original image* is reached.

This implementation demonstrated that by constructing a Laplacian pyramid for the spatial filter finer motion detail would be revealed, whereas the color variation, the property to be analyzed, was less evident.

4.2.4 Performance optimized EvmGdownIIR

Because the OpenCV Java binding was not complete at the time, the Java desktop implementations were not executing fast enough for real-time processing. Thus, a C/C++ implementation of the method that demonstrated better color variation was developed. In addition, after studying the performance of the application, detailed on section 5.2, the resize operations proved to be computationally expensive. Therefore, the *Resize down* and *Resize up* steps were modified to faster resize operations that did not alter the resulting image.

Resize down

This step was changed to a single resize operation using the OpenCV interpolation method named *area*, which produces a similar result to the one provided by using a Gaussian filter

and downsampling the image. However, instead of, iteratively downsampling the frame multiple times, the resize in this step is done in only once to a predefined size.

Resize up

This step was also modified to a single resize operation using the linear interpolation method, which produces a good result in short amount of time.

4.3 Face detection

The *Face detector* module uses the OpenCV object detector initially proposed on [VJ01] and improved on [LM02]. Which has been previously trained to detect faces.

Because object detectors are computationally expensive, therefore, in order to improve performance, a minimum size for the face detector was set to 40% of the frame width and height. In addition, since it was expected for the person to remain still during the reading, the face detector was set to execute only once a second.

Face tracking is simply done by matching the previous and newly detected faces to the closest one.

- If there are less newly detected faces than the previously detected ones, then match each new face to the nearest older face. Any older face that is not matched with a newly detected face is marked for deletion. However, it is only deleted if it fails to find a match on the next time the face detector executes. This measure allows the face detector to miss a detection of a face one time.
- Otherwise, if the number of newly detected faces is equal or more than the older faces, then match each older face to the nearest newly detected face. Any newly detected face that is not matched with an older face is then marked as a new face.

Another performance boost was achieved by only magnifying the each face rectangle instead of the whole frame. Because of this, the face rectangle must remain as still as possible, in order to introduce as few artifacts and noise to the Eulerian Video Magnification method. To achieve this, the position and size of the face's rectangle that is fed into the Eulerian Video Magnification method is interpolated between the previous and newly detected faces, if the distance between the two, d , is less than one third of the previous face's rectangle width, w :

$$d < \frac{w}{3} \quad (4.4)$$

And, the interpolation percentage, r , used is the ratio between these values:

$$r = \frac{3d}{w} \quad (4.5)$$

Implementation details

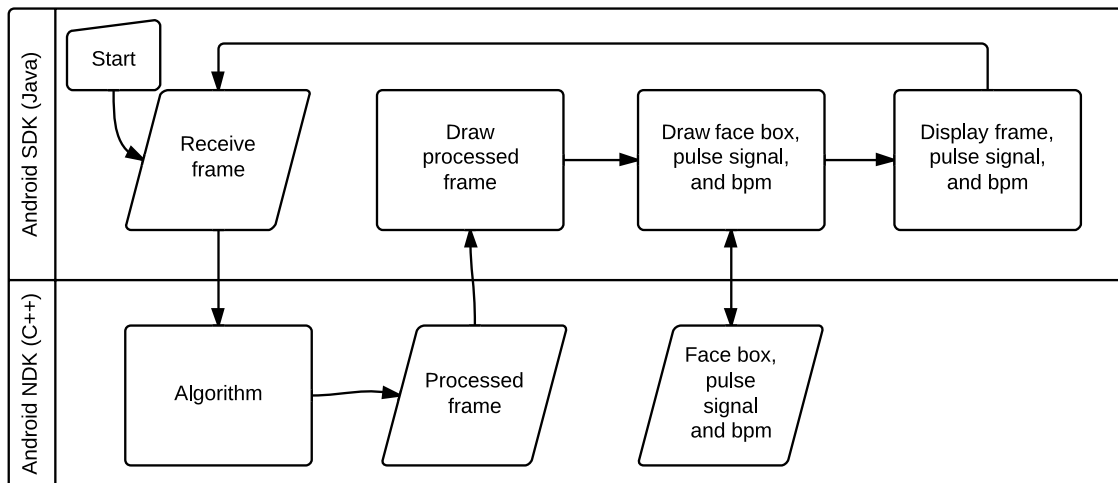


Figure 4.4: Integration workflow between Android native and Java parts.

4.4 Signal validations

...

4.5 Heart rate estimation

...

4.6 Performance optimizations

...

4.7 Android integration

...

4.8 Chapter summary

...

Implementation details

Chapter 5

Results

5.1 Heart rate comparison

5.2 Algorithm performance

5.3 Chapter summary

Results

Chapter 6

Conclusions

6.1 Objective satisfaction

6.2 Future work

Conclusions

References

- [BA83] P. Burt and E. Adelson. The laplacian pyramid as a compact image code. *Communications, IEEE Transactions on*, 31(4):532–540, 1983.
- [Com94] P. Comon. Independent component analysis, a new concept? *Signal processing*, 36(3):287–314, 1994.
- [Gre97] EF Grenaker. Radar sensing of heartbeat and respiration at a distance with applications of the technology. In *Radar 97 (Conf. Publ. No. 449)*, pages 150–154. IET, 1997.
- [GSMP07] M. Garbey, N. Sun, A. Merla, and I. Pavlidis. Contact-free measurement of cardiac pulse based on the analysis of thermal imagery. *Biomedical Engineering, IEEE Transactions on*, 54(8):1418–1426, 2007.
- [IDC13] IDC. Android marks fourth anniversary since launch with 75.0 in third quarter, according to idc. <http://www.idc.com/getdoc.jsp?containerId=prUS23771812>, January 2013.
- [Its13] OpenCV Developers Team: Itseez. About | opencv. <http://opencv.org/about.html>, January 2013.
- [LFSK06] C. Liu, W.T. Freeman, R. Szeliski, and S.B. Kang. Noise estimation from a single image. In *Computer Vision and Pattern Recognition, 2006 IEEE Computer Society Conference on*, volume 1, pages 901–908. IEEE, 2006.
- [LM02] Rainer Lienhart and Jochen Maydt. An extended set of haar-like features for rapid object detection. In *Image Processing. 2002. Proceedings. 2002 International Conference on*, volume 1, pages I–900. IEEE, 2002.
- [LOCS95] Daniel A Litvack, Tim F Oberlander, Laurel H Carney, and J Philip Saul. Time and frequency domain methods for heart rate variability analysis: a methodological comparison. *Psychophysiology*, 32(5):492–504, 1995.
- [LTF⁺05] C. Liu, A. Torralba, W.T. Freeman, F. Durand, and E.H. Adelson. Motion magnification. In *ACM Transactions on Graphics (TOG)*, volume 24, pages 519–526. ACM, 2005.
- [NI10] Bistra Nenova and Ivo Iliev. An automated algorithm for fast pulse wave detection. *International Journal Bioautomation*, 14(3):203–216, 2010.
- [PB90] Stephen W Porges and Robert E Bohrer. The analysis of periodic processes in psychophysiological research. 1990.

REFERENCES

- [Phi13] Philips. Philips vital signs camera. <http://www.vitalsignscamera.com>, January 2013.
- [PMP10] M.Z. Poh, D.J. McDuff, and R.W. Picard. Non-contact, automated cardiac pulse measurements using video imaging and blind source separation. *Optics Express*, 18(10):10762–10774, 2010.
- [PMP11] M.Z. Poh, D.J. McDuff, and R.W. Picard. Advancements in noncontact, multiparameter physiological measurements using a webcam. *Biomedical Engineering, IEEE Transactions on*, 58(1):7–11, 2011.
- [Por13] Fraunhofer Portugal. Fraunhofer portugal - about us. <http://www.fraunhofer.pt>, January 2013.
- [Tec13] ViTrox Technologies. What’s my heart rate. <http://www.whatsmyheartrate.com>, May 2013.
- [TRaK02] Mika P Tarvainen, Perttu O Ranta-aho, and Pasi A Karjalainen. An advanced detrending method with application to hrv analysis. *Biomedical Engineering, IEEE Transactions on*, 49(2):172–175, 2002.
- [UT93] S.S. Ulyanov and V.V. Tuchin. Pulse-wave monitoring by means of focused laser beams scattered by skin surface and membranes. In *OE/LASE’93: Optics, Electro-Optics, & Laser Applications in Science & Engineering*, pages 160–167. International Society for Optics and Photonics, 1993.
- [VJ01] Paul Viola and Michael Jones. Rapid object detection using a boosted cascade of simple features. In *Computer Vision and Pattern Recognition, 2001. CVPR 2001. Proceedings of the 2001 IEEE Computer Society Conference on*, volume 1, pages I–511. IEEE, 2001.
- [VSN08] Wim Verkruyse, Lars O Svaasand, and J Stuart Nelson. Remote plethysmographic imaging using ambient light. *Optics express*, 16(26):21434–21445, 2008.
- [WRS⁺12] Hao-Yu Wu, Michael Rubinstein, Eugene Shih, John Guttag, Frédo Durand, and William T. Freeman. Eulerian video magnification for revealing subtle changes in the world. *ACM Trans. Graph. (Proceedings SIGGRAPH 2012)*, 31(4), 2012.